

Simulating Spectral Sampling of the MICA Files with Enfys: InAs SNR10x MWIR

August 15, 2025

1 Abstract

Here we investigate the expected ability of a new remote-sensing stand-off infrared spectrometer, *Enfys*, for the ExoMars Rosalind Franklin Rover, to resolve diagnostic features of the reflectance spectra of key minerals previously identified on the surface of Mars.

In this notebook we simulate Enfys performance using an InAs sensor, that is subject to thermal noise from the Enfys structure.

Here we imagine a worst case scenario where an SNR of 1/10 the best case is achieved in the InAs channel.

2 Introduction

As discussed in [previous notebooks](#), Enfys is an Infra-Red point-spectrometer for the ExoMars Rosalind Franklin rover.

In these notebooks we explore how we expect Enfys to sample minerals of interest on the Mars surface.

In this notebook, we consider the expected performance of Enfys using an InAs Mid-Wave Infrared sensor for the wavelength range of 1.6 - 3.1 μm .

Included in the modelling is the noise expected for the sensor, as provided by the EXM-EN-MTX-ABU-0001_Enfys_Science_Traceability_Matrix_3-0 document.

2.1 Problem

This optical design is functionally different to the Acousto-Optic Tunable Filter design of ISEM, so it is important that the design is demonstrated to fulfil the performance requirements of the instrument, and that expectations of performance can be set prior to laboratory tests of the completed instrument.

At time of writing, the spectrometer is in the development stage. This notebook describes how the Spectral Parameters Toolkit¹ can be used to convert the developing expectations of the spectral range and resolution and radiometric sensitivity of the instrument into predictions of its ability to discriminate mineral species and phases on the Mars surface at the rover landing site.

¹Stabbins, R., & Grindrod, P. (2024). The Spectral Parameters Toolkit (sptk): Release v0.1 (Version v0.1) [Computer software]. Zenodo. <https://doi.org/10.5281/zenodo.10694286>

3 Method

To investigate the expected performance we use high-resolution laboratory measured spectra of a range of minerals expected at the Mars surface, and computationally simulate the sampling of the spectra with the expected spectral response of *Enfys*. We will perform repeat sampling with synthetically added noise to investigate the spectral and radiometric limits of the instrument.

We will produce a number of visualisations to investigate the predicted instrument performance, and demonstrate the limits of the instrument discriminatory ability.

3.1 Spectral Resampling

We take a simple approach to resampling, as follows.

We neglect reflectance calibration uncertainty, and any systematic uncertainty when merging the data from the two separate photodiodes, and assume that the linear-variable filter movement is perfectly calibrated such that a measurement at centre-wavelength λ_{cwl} can be requested with arbitrary precision.

We assume an observation of λ_{cwl} has a Cauchy transmission profile, and, prior to component characterisation, a spectral resolving power of $\lambda/\Delta\lambda = 100$ is achieved for all λ_{cwl} , such that the Cauchy transmission profile for λ_{cwl} has a full-width-at-half-maximum of $\Delta\lambda = \lambda_{cwl}/100$.

The Cauchy transmission profile $T(\lambda|\lambda_{cwl})$ is given by

$$T(\lambda|\lambda_{cwl}) = \frac{\gamma^2}{(\lambda - \lambda_{cwl})^2 + \gamma^2}$$

where $\gamma = \Delta\lambda/2 = (\lambda_{cwl}/100)/2$.

We assume that a high-resolution laboratory spectrum approximates the continuous function $R(\lambda)$, such that the calibrated reflectance at λ_{cwl} has the expected value $R[\lambda_{cwl}]$ of:

$$R[\lambda_{cwl}] = \frac{\int R(\lambda)T(\lambda|\lambda_{cwl})d\lambda}{\int T(\lambda|\lambda_{cwl})d\lambda}$$

To simulate noise we add ε , an additive noise term that we draw from the Gaussian distribution with variance:

$$\sigma^2 = \frac{\sqrt{R[\lambda_{cwl}]}}{\text{SNR}}$$

under the assumption that maximum SNR is only achieved for $R = 1$, and the noise is dominated by counting noise, such that σ follows a Poissonian relationship with the signal R .

For each observation $R[\lambda_{cwl}]$ we draw 100 repeat samples, such that

$$R_i[\lambda_{cwl}] = R[\lambda_{cwl}] + \varepsilon_i$$

with

$$\varepsilon_i \sim \mathcal{N} \left(0, \frac{\sqrt{R[\lambda_{cwl}]}}{\text{SNR}} \right)$$

This gives an indication of the expected noise of a calibrated *Enfys* acquired spectrum.

3.2 Analysis

For now, we simply produce plots of the resampled spectra, and visually inspect the plots to consider the ability to resolve diagnostic spectral features against the noise. We will include more details of analysis as that is developed here.

3.3 Notebook and Environment Setup

Here we include the required code blocks for setting up this notebook for this study.

```
[1]: %load_ext autoreload
      %autoreload 2
      %config InlineBackend.figure_format = 'svg' # note - requires Inkscape to be
      ↪ installed, and to be accessible from the terminal
      # on macos, add inkscape to the path with `sudo ln -s /Applications/Inkscape.
      ↪ app/Contents/MacOS/inkscape /usr/local/bin/inkscape`
      %matplotlib inline
```

We will be using the Spectral Parameters Toolkit. First we import the required modules.

```
[2]: import sptk.config as cfg
      from sptk.material_collection import MaterialCollection
      from sptk.instrument import Instrument
      from sptk.observation import Observation
      from sptk.spectral_library_analyser import SpectralLibraryAnalyser as sla
```

We set our simulation resolution range to extend the range of *Enfys*, of 0.9 - 3.1 μm , such that the long tails of the Cauchy distribution of the linear-variable-filter profiles can be accommodated during sampling. We leave our spectral resolution at 1 nm for all λ .

```
[3]: cfg.update_sample_res('wvl_min', 380)
      cfg.update_sample_res('wvl_max', 3300)
```

Finally we set the project name.

```
[4]: PROJECT_NAME = 'enfys_mica_inas_mwir_snr10x'
```

4 Data

For this study we are using a spectral library of minerals identified through CRISM analysis ([the MICA files](#))². Here we use the laboratory spectra that have been identified that best represent the

²Viviano-Beck, C. E., Seelos, F. P., Murchie, S. L., Kahn, E. G., Seelos, K. D., Taylor, H. W., Taylor, K., Ehlmann, B. L., Wiseman, S. M., Mustard, J. F., & Morgan, M. F. (2014). Revised CRISM spectral parameters and summary

type localities presented in the MICA files. These can be accessed [here](#), and are sourced from the RELAB and USGS public spectral libraries.

We have grouped the spectral library into categories as follows, and included this categorisation in the `sptk` distribution, accessed from the `sptk.material_sets` module.

```
[5]: from sptk.material_sets import MICA_SET
```

We use the `sptk.material_collection` module to parse the spectral library:

```
[6]: matcol = MaterialCollection(
    MICA_SET,
    'MICA_new_dirtree',
    PROJECT_NAME,
    load_existing=False,
    balance_classes=False,
    random_bias_seed=None,
    allow_out_of_bounds=True,
    plot_profiles=False,
    export_df=True)
```

Building new enfys_mica_inas_mwir_snr10x MaterialCollection DF

Category

species|subgroup|group|
data_id status

Iron Oxides & Primary Silicates

-clinopyroxene|pyroxenes|inosilicates|
-NBPP22.csv loaded
-orthopyroxene|pyroxenes|inosilicates|
-LAPP47B.csv loaded
-fayalite|olivines|nesosilicates|
-LAP005.csv loaded
-forsterite|olivines|nesosilicates|
-LASC02.csv loaded
-hematite|hematites|oxides|
-GDS27.csv loaded
-plagioclase|feldspars|tectosilicates|
-NCLS04.csv loaded

Sulfates

-alunite|hydrous_sulphates|sulphates|
-F1CC08B.csv loaded
-bassanite|hydrous_sulphates|sulphates|
-GDS145.csv loaded

products based on the currently detected mineral diversity on Mars. Journal of Geophysical Research: Planets, 119(6), 1403–1431. <https://doi.org/10.1002/2014JE004627>

-gypsum|hydrous_sulphates|sulphates|
-LASF41A.csv loaded
-jarosite|hydrous_sulphates|sulphates|
-LASF21A.csv loaded
-kieserite|hydrous_sulphates|sulphates|
-F1CC15.csv loaded
-mg-sulphate|hydrous_sulphates|sulphates|
-799F366.csv loaded

Phyllosilicates

-chlorite|chlorites|phyllosilicates|
-LACL14.csv loaded
-kaolinite|kaolinites|phyllosilicates|
-LAKA04.csv loaded
-illite|micas|phyllosilicates|
-LAIL02.csv loaded
-margarite|micas|phyllosilicates|
-GDS106.csv loaded
-vermiculite|micas|phyllosilicates|
-c1cy29.csv loaded
-serpentine|serpentine|phyllosilicates|
-LALZ01.csv loaded
-montmorillonite|smectites|phyllosilicates|
-LAM002.csv loaded
-nontronite|smectites|phyllosilicates|
-NCJB26.csv loaded
-saponite|smectites|phyllosilicates|
-LASA58.csv loaded
-talc|talcs|phyllosilicates|
-GDS23.csv loaded

Carbonates

-calcite|calcites|carbonates|
-CAGR01.csv loaded
-magnesite|calcites|carbonates|
-F1CC06B.csv loaded

Hydrated Silicates & Halides

-halite|chlorides|halides|
-HS433.csv loaded
-prehnite|prehnites|inosilicates|
-LAZE03.csv loaded
-siliceous-sinter|hydrated-silicas|mineraloids|
-bkr1jb329c.csv loaded
-epidote|epidotes|sorosilicates|
-GDS26.csv loaded
-analcime|zeolites|tectosilicates|
-GDS1.csv loaded

Loading 29 of entries complete.

Exporting the Material Collection to CSV and Pickle formats...

We can use the `material_collection` plotting routines to visualise the library.

```
[7]: ax = matcol.plot_profiles(stacked=True, scope='categories', groupby='Species')
```

Categories grouped by Species




Figure 2 High resolution MICA Files spectral library, arranged by category and grouped by mineral

species.

We can see that not all of the entries cover the complete 0.9 - 3.1 μm range of *Enfys*, so in future work we should look to replace these MICA files laboratory type spectra with spectra that all span the entire range.

We can also produce a waterfall visualisation, an alternative plotting method that presents the information in a condensed format, that is useful for highlighting where absorption features lie.

```
[8]: axes = matcol.plot_profiles(waterfall=True, scope='all', groupby='Category')
```




Figure 3 Waterfall plot of the high resolution MICA Files spectral library, arranged by category and grouped by mineral species.

A more useful version of this plot format is to show band locations after continuum removal. We can use the `SpectralLibraryAnalyser.remove_continuum()` function to produce a duplicate of the `MaterialCollection` but with the data continuum-removed.

```
[9]: from sptk.spectral_library_analyser import SpectralLibraryAnalyser as sla

matcol_sla = sla(matcol)
matcol_cr = matcol_sla.remove_continuum()
```

We can plot the profiles to illustrate the continuum removal. Here we use the `scope = 'all'` version of the profile plot.


```
[10]: axes = matcol_cr.plot_profiles(stacked=True, scope='all', groupby='Category')
```




Figure 4 Continuum-removed high resolution MICA Files spectral library, grouped by Category.

Now we plot the waterfall representation:

```
[11]: axes = matcol_cr.plot_profiles(waterfall=True, scope='all', groupby='Category')
```




Figure 5 Continuum-removed waterfall plot of high resolution MICA Files spectral library, grouped by Category.

Again, we can see that the waterfall plot represents the band depth information of the `MaterialCollection` in a compact format. Scanning vertically allows for quick identification of common band locations, their depths, widths, and asymmetries.

5 Results

5.1 Simulating the EnfyS Spectral Response

To simulate EnfyS, rather than providing a file with the expected spectral response of each position of the Linear-Variable Filter (LVF, above), we specify the spectral range and resolution and sampling regime, and generate a set of central-wavelengths and bandwidths. We do this with the `InstrumentBuilder` class of the `sptk.instrument` module.

```
[12]: from sptk.instrument import InstrumentBuilder
```

The expected spectral resolution of Enfys is $\lambda/\Delta\lambda = 100$, and the range when using the InGa sensor is 0.9 - 3.1 μm .

5.1.1 Signal-to-Noise Ratio

We obtain the expected SNR from the enfys_inas_mwir_snr.csv file, derived from the EXM-EN-MTX-ABU-0001_Enfys_Science_Traceability_Matrix_3-0 spreadsheet.

We have adapted the InstrumentBuilder function to accept a pandas Series mapping wavelength to SNR. The InstrumentBuilder then performs interpolation to the simulated wavelength range.

We read in the SNR file:

```
[13]: import pandas as pd

      enfys_snr = pd.read_csv('../data/instruments/enfys_snr.csv', index_col=0)
```

and get the SNR data for the InAs MWIR sensor,

```
[14]: inas_mwir_snr = enfys_snr['InAs MWIR SNR']
```

and the InGaAs SWIR sensor,

```
[15]: ingaas_swir_snr = enfys_snr['InGaAs SWIR SNR']
```

Here we are going to artificially scale the InAs SNR such that the max SNR is 50.

```
[16]: import numpy as np
      inas_mwir_snr_max = np.nanmax(inas_mwir_snr.to_numpy())
      inas_mwir_snr_max
```


```
[16]: 674.37153190000004
```

```
[17]: inas_mwir_snr = inas_mwir_snr / 10 ## 50 / inas_mwir_snr_max
```

and then we splice these together, using the higher value SNR where the wavelength ranges overlap:

```
[18]: inas_mwir_snr.plot(logy=True)
      ingaas_swir_snr.plot()
```

```
[18]: <AxesSubplot: xlabel='Wavelength(um)'\>
```



```
[19]: # concatenate the inas_mwir_snr and the ingaas_swir, using the highest SNR for
      overlapping wavelength values
      enfys_inas_mwir_snr = ingaas_swir_snr.combine_first(inas_mwir_snr)
```

We also convert the wavelengths from μm to nm:

```
[20]: enfys_inas_mwir_snr.index = enfys_inas_mwir_snr.index * 1000
```


for computational convenience, we should set values of $\text{SNR} \leq 1$ to NaN, to interpret these as ‘bad bands’

```
[21]: # where SNR <=1 to NaN, to interpret these as 'bad bands'
      import numpy as np
      enfys_inas_mwir_snr = enfys_inas_mwir_snr.where(enfys_inas_mwir_snr > 3,
      other=np.nan)
```

```
[22]: enfys_inas_mwir_snr.index.name = 'wavelength'
```

```
[23]: g = enfys_inas_mwir_snr.plot(logy=True)
      # set the x and y axes
      g.set_xlabel('Wavelength (nm)')
      g.set_ylabel('SNR')
```

```
[23]: Text(0, 0.5, 'SNR')
```



5.1.2 Spectral Resolving Power


We can also load in new data for the spectral resolving power of the SWIR and MWIR channels of Enfys.

This data has been digitised from the plot reported in the Enfys schedule review 20250807 ppt, by Matt Gunn.

```
[24]: enfys_swir_respow = pd.read_csv('../data/instruments/enfys_swir_respow.csv',  
    ↪ index_col=0)  
enfys_mwir_respow = pd.read_csv('../data/instruments/enfys_mwir_respow.csv',  
    ↪ index_col=0)
```

```
[25]: g=enfys_swir_respow.plot(label='SWIR', marker='o')  
enfys_mwir_respow.plot(marker='o',label='MWIR', xlabel='Wavelength (nm)',  
    ↪ ylabel=r'$\lambda/\Delta\lambda$ Resolving Power',ax=g)  
g.legend(['SWIR', 'MWIR'])
```

```
[25]: <matplotlib.legend.Legend at 0x30e297d00>
```



We now need to merge these datasets, following the same rule as the SNR merger, i.e. ensuring that the correct channel is used for the overlapping data.

```
[26]: # get the max wavelength where ingaas_swir_snr is not nan
ingaas_swir_snr = ingaas_swir_snr[~np.isnan(ingaas_swir_snr)]
max_swir_wavelength = ingaas_swir_snr.index[-1] * 1000 # convert from μm to nm
max_swir_wavelength
```

[26]: 1650.0

```
[27]: # get the enfys_swir_respow data where the wavelength is less than the
      ↳ max_swir_wavelength
enfys_swir_respow = enfys_swir_respow[enfys_swir_respow.index <
      ↳ max_swir_wavelength]
enfys_swir_respow
```

```
[27]:           Wavelength    Resolving Power
114.172    952.350
100.267    984.300
 99.765   1016.886
```

1050.340	100.625
1084.064	103.015
1117.569	106.020
1151.685	104.326
1186.342	110.645
1221.214	113.238
1256.182	116.001
1291.127	117.274
1326.348	120.119
1361.845	119.359
1397.410	125.192
1432.961	125.249
1468.431	127.087
1503.552	129.283
1538.147	129.767
1572.489	130.274
1606.157	131.627
1638.722	130.653

```
[28]: # get the enfys_mwir_respow where the wavelength is greater than the
      ↪ max_swir_wavelength
enfys_mwir_respow = enfys_mwir_respow[enfys_mwir_respow.index >
      ↪ max_swir_wavelength]
enfys_mwir_respow
```


```
[28]:
```

Wavelength	Resolving Power
1673.354	100.000
1732.701	101.075
1792.699	102.803
1854.283	105.507
1916.113	109.126
1978.328	111.035
2041.291	114.507
2104.575	117.912
2232.792	120.042
2298.070	121.352
2363.760	124.419

```
[29]: # now merge the datasets
enfys_respow = pd.concat([enfys_swir_respow, enfys_mwir_respow])
```

```
[30]: g=enfys_respow.plot(xlabel='Wavelength (nm)', ylabel=r'$\lambda/\Delta\lambda$,
      ↪ Spectral Resolving Power')
      g.legend(['SWIR & MWIR'])
```

```
[30]: <matplotlib.legend.Legend at 0x30d781a20>
```

We now need to pass this data to the InstrumentBuilder function.

5.1.3 Instrument Build

Now we will pass this to the InstrumentBuilder:

```
[31]: enfys_builder = InstrumentBuilder(
    instrument_name='enfys_inas_mwir',
    instrument_type='lvf',
    sampling='nyquist',
    resolution = enfys_respow,
    spectral_range=[900, 3101],
    snr = enfys_inas_mwir_snr)
```

Exporting instrument to ../data/instruments/enfys_inas_mwir.csv...

```
[32]: from sptk.instrument import Instrument
```

Now we build the transmission profiles for each LVF position, according to the list of CWLs and FWHMs generated above. These are accessed by reading the .csv file produced by the InstrumentBuilder procedure.

When generating the profiles, we turn the CWL and FWHM into a profile by assuming a distribution

function. For Enfys we use the Cauchy distribution, which behaves similarly to a Gaussian, but with wider ‘wings’. It is parameterised by the full-width-at-half-maximum, as the distribution strictly has no finite variance. We have implemented a `build_cauchy_filter` method for the `Instrument` class.

```
[33]: enfys = Instrument(  
    name = 'enfys_inas_mwir', # filename of the instrument information,   
    ↪held in /software/data/instruments/  
    project_name = matcol.project_name, # the project directory name,   
    ↪hosting outputs in the /spectral_parameter_studies/ directory  
    shape = 'cauchy',  
    load_existing=False, # if True, load existing instrument data from   
    ↪the project directory  
    plot_profiles=False, # if true, plot the transmission profiles   
    ↪during construction  
    export_df=True) # if True, export the instrument transmission   
    ↪profiles to the project directory
```

Building Instrument...

Building new DataFrame for 'enfys_inas_mwir'...

Exporting the Instrument to CSV and Pickle formats...

```
[34]: fig ,ax = enfys.plot_instrument_characteristics()
```

Plotting Instrument Transmission...

/opt/homebrew/Caskroom/miniconda/base/envs/enfys_sptk/lib/python3.10/site-packages/seaborn/_oldcore.py:200: FutureWarning: elementwise comparison failed; returning scalar instead, but in the future will perform elementwise comparison
if palette in QUAL_PALETTES:

Plotting Maximum Signal-to-Noise Ratio...

Plotting FWHM...

Plotting Spectral Resolution...

Plots exported to ../projects/enfys_mica_inas_mwir_snr10x/instrument




Figure 6 Simulated spectral response and resolution of *Enfys*, based on preliminary expectations of the spectral and radiometric response.

The data shown in figure 3 approximates the expected performance of *Enfys*. Once the instrument is calibrated at the component and assembled levels, this data can be replaced with empirically acquired data.

5.2 Sampling MICA with Enfys

We now make an `Observation` object by sampling the MICA files `MaterialCollection` with the `Enfys Instrument`.

```
[35]: from sptk.observation import Observation
```

```
[36]: obs = Observation(
        material_collection=matcol,
```

```
instrument = enfys,  
load_existing=False,  
plot_profiles=False,  
export_df=True)
```

Building new Observation DataFrame
for 'enfys_mica_inas_mwir_snr10x'...
Sampling the Material Collection with the Instrument...
Sampling complete.
Exporting the Observation Pickle format...
Observation export complete.

We can plot the sampled spectra against the input high-spectra below:

```
[37]: axes = obs.plot_profiles(scope='categories', groupby='Species', stacked=True, ↵  
    ↪ hires_under=True)
```




Figure 7 The MICA files as sampled by *Enfys*, with the original high-resolution spectra plotted

underneath each.

We can now see where there are deviations of the *Enfys* sampled spectra from the high-resolution source spectra.

For example, illite, plotted below for clarity.

```
[38]: axes = obs.plot_profiles(scope='illite', groupby='Sample ID', stacked=False, ↵  
    ↪ hires_under=True)
```

Enfys_Inas_Mwir Sampled Illite grouped by Sample Id




Figure 8 Deviation of the *Enfys* sampled spectrum from the high-resolution input spectrum for Illite.

We note that there is an effect by which peaks and troughs in the spectra are reduced. This is due to the long tails of the Cauchy filter profiles. Despite the reduction of absolute band depth, the features of Illite (fig. 8) are still resolvable.

We can produce a continuum-removed waterfall plot to check that the major bands are resolved.

```
[39]: from sptk.spectral_library_analyser import SpectralLibraryAnalyser as sla  
  
obs_cr = sla(obs).remove_continuum()
```

```
[40]: axes = obs_cr.plot_profiles(waterfall=True, scope='all', groupby='Category')
```




Figure 9 Continuum-Removed waterfall plot of the MICA files as sampled by *Enfys*.

Comparing Figure 9 to Figure 5, we can see that the weaker sub-2 μm bands are less well resolved.

We can save this data

```
[41]: obs_cr.export_main_df()
```

Exporting the Observation Pickle format...
Observation export complete.

5.3 Adding Synthetic Noise

We have the capability to add synthetic noise to the dataset, by redrawing a sample from the dataset with noise added according to the signal-to-noise ratio specified for the given channel.

We can optionally redraw multiple noisy samples for each entry, to give an indication of the statistical uncertainty, or for example draw a single noisy sample, to produce a single-observation mission-realistic dataset.

For an entry with reflectance $R[\lambda_{cwl}]$, we redraw a sample:

$$\tilde{R}_i[\lambda_{cwl}] = R[\lambda_{cwl}] + \epsilon_i$$

where ϵ_i is drawn from:

$$\epsilon_i \sim \mathcal{N}(0, \sigma[\lambda_{cwl}]^2)$$

where $\sigma[\lambda_{cwl}]$ is the noise on the reflectance measurement, given by:

$$\sigma[\lambda_{cwl}] = R[\lambda_{cwl}] / \text{SNR}[\lambda_{cwl}]$$

First, let's draw a single noisy entry:

```
[42]: noise_obs = obs.add_noise(1, seed=0)
```

```
[43]: axes = obs.plot_profiles(scope='illite', groupby='Sample ID', stacked=False,
    ↪ hires_under=True, ci=False)
```

Enfys_Inas_Mwir Sampled Illite grouped by Sample Id with Noise




Figure 10 Noisy resampled example spectrum for Illite, showing 1 noisy sample according to the spectral signal-to-noise ratio illustrated in figure 3.B.

We can also plot the mean observed line, with shading indicating the standard deviation of the noisy data, shown below in figure 7, where the 'ci' keyword has been applied ('ci': confidence-interval).

```
[44]: axes = obs.plot_profiles(scope='illite', groupby='Sample ID', stacked=False,
    ↪ hires_under=True, ci=True)
```


Enfys_Inas_Mwir Sampled Illite grouped by Sample Id with Noise




Figure 11 Mean of Noisy resampled example spectrum for Illite, showing mean and standard deviation of 100 repeat samples according to the spectral signal-to-noise ratio illustrated in figure 3.B.

Here we find that the standard deviation is almost negligible at the scale of the figure.

Plotting out all of the groups, we again see that the noise, at this scale, is indeed negligible for most entries. Here we opt to plot the complete range of noisy samples, rather than averaging and showing the (very small) standard deviation.

```
[45]: axes = obs.plot_profiles(scope='categories', groupby='Species', stacked=True,
    ↪ hires_under=True, ci=False)
```




Figure 12 The MICA files as sampled by *Enfys*, with repeat samples drawn according to the ex-







pected spectral Signal-to-Noise Ratio, with the original high-resolution spectra plotted underneath each.

The key conclusion to draw from the above figure is that expected noise for *Enfys* does not significantly intrude on the diagnostic features of these type specimen of minerals identified on the Mars surface.


We can illustrate this in greater detail by plotting each entry on a separate figure, with the y-axis clipped to the spectral range of entry. Note that for sample with a small spectral range (i.e. samples that are spectrally flat), this stretching will appear to amplify noise, so the y-axis range must be considered when assessing the noise.

```
[46]: axes = obs.plot_profiles(scope='samples', groupby='Species', stacked=False,
    ↪ hires_under=True, ci=False)
```


Enfys_Inas_Mwir Sampled Samples grouped by Species with Noise (1/5)









Enfys_Inas_Mwir Sampled Samples grouped by Species with Noise (2/5)



Enfys_Inas_Mwir Sampled Samples grouped by Species with Noise (3/5)



Enfys_Inas_Mwir Sampled Samples grouped by Species with Noise (4/5)






Figure 13 Detailed view of each entry in the spectral library, showing the scale of noise introduced,

relative to the features of each spectrum, and the reflectance range.

We can export the dataset:

```
[47]: obs.export_main_df()
```

```
Exporting the Observation Pickle format...  
Observation export complete.
```

5.4 Continuum Removal

We can apply continuum removal, even to the noisy data.

```
[48]: from sptk.spectral_library_analyser import SpectralLibraryAnalyser as sla  
  
cr_tool = sla(obs)  
cr_obs = cr_tool.remove_continuum()
```

```
[49]: axes = cr_obs.plot_profiles(scope='categories', groupby='Species',  
    ↪stacked=True, hires_under=False, ci=False)
```




Figure 14 Continuum-Removed entries of the MICA Files, as sampled by *Enfys*, including syn-

thetic noise.

Again, we can export this:

```
[50]: cr_obs.export_main_df()
```

```
Exporting the Observation Pickle format...
Observation export complete.
```

5.5 Waterfall Visualisation

For waterfall visualisations, we have the advantage of being able to visualise the noise introduced for each of multiple entries in 2D.

Each bar of the visualisation is composed of a series of rows, with each row as a redrawn noisy sample.

For this, we redraw with 128 repeat samples.

```
[51]: obs.add_noise(128, seed=0)
```

```
[51]:
```

	Category	Root Data ID	Sample ID	\
Data ID				
CAGR01 Calcite.	carbonates	CAGR01 Calcite	CAGR01	
CAGR01 Calcite.1	carbonates	CAGR01 Calcite	CAGR01	
CAGR01 Calcite.10	carbonates	CAGR01 Calcite	CAGR01	
CAGR01 Calcite.100	carbonates	CAGR01 Calcite	CAGR01	
CAGR01 Calcite.101	carbonates	CAGR01 Calcite	CAGR01	
...	...	...	...	
799F366 Mg-sulphate.95	sulfates	799F366 Mg-sulphate	799F366	
799F366 Mg-sulphate.96	sulfates	799F366 Mg-sulphate	799F366	
799F366 Mg-sulphate.97	sulfates	799F366 Mg-sulphate	799F366	
799F366 Mg-sulphate.98	sulfates	799F366 Mg-sulphate	799F366	
799F366 Mg-sulphate.99	sulfates	799F366 Mg-sulphate	799F366	

	Species	Subgroup	Group	\
Data ID				
CAGR01 Calcite.	calcite	calcites	carbonates	
CAGR01 Calcite.1	calcite	calcites	carbonates	
CAGR01 Calcite.10	calcite	calcites	carbonates	
CAGR01 Calcite.100	calcite	calcites	carbonates	
CAGR01 Calcite.101	calcite	calcites	carbonates	
...	...	...	...	
799F366 Mg-sulphate.95	mg-sulphate	hydrous_sulphates	sulphates	
799F366 Mg-sulphate.96	mg-sulphate	hydrous_sulphates	sulphates	
799F366 Mg-sulphate.97	mg-sulphate	hydrous_sulphates	sulphates	
799F366 Mg-sulphate.98	mg-sulphate	hydrous_sulphates	sulphates	
799F366 Mg-sulphate.99	mg-sulphate	hydrous_sulphates	sulphates	

Library Sample Description Date Added \

Data ID				
CAGR01 Calcite.	MICA_new_dirtree		NaN	NaN
CAGR01 Calcite.1	MICA_new_dirtree		NaN	NaN
CAGR01 Calcite.10	MICA_new_dirtree		NaN	NaN
CAGR01 Calcite.100	MICA_new_dirtree		NaN	NaN
CAGR01 Calcite.101	MICA_new_dirtree		NaN	NaN
...	...	...	...	...
799F366 Mg-sulphate.95	MICA_new_dirtree		NaN	NaN
799F366 Mg-sulphate.96	MICA_new_dirtree		NaN	NaN
799F366 Mg-sulphate.97	MICA_new_dirtree		NaN	NaN
799F366 Mg-sulphate.98	MICA_new_dirtree		NaN	NaN
799F366 Mg-sulphate.99	MICA_new_dirtree		NaN	NaN

Viewing Geometry ... 3011.83736494 3021.57325265 \

Data ID	...			
CAGR01 Calcite.	NaN	...	0.622482	0.579071
CAGR01 Calcite.1	NaN	...	0.599965	0.743005
CAGR01 Calcite.10	NaN	...	0.584973	0.652090
CAGR01 Calcite.100	NaN	...	0.710601	0.633514
CAGR01 Calcite.101	NaN	...	0.715903	0.621524
...	...	...	...	...
799F366 Mg-sulphate.95	NaN	...	0.045230	0.046224
799F366 Mg-sulphate.96	NaN	...	0.046699	0.049784
799F366 Mg-sulphate.97	NaN	...	0.042550	0.046324
799F366 Mg-sulphate.98	NaN	...	0.044031	0.041885
799F366 Mg-sulphate.99	NaN	...	0.039099	0.044458

3031.3119922 3041.05356775 3050.79796359 3060.54516415 \

Data ID				
CAGR01 Calcite.	0.583521	0.659142	0.594213	0.515316
CAGR01 Calcite.1	0.686747	0.638356	0.733249	0.562226
CAGR01 Calcite.10	0.649290	0.643655	0.695801	0.625918
CAGR01 Calcite.100	0.691883	0.601379	0.609805	0.738102
CAGR01 Calcite.101	0.678623	0.691925	0.655394	0.661303
...	...	...	...	...
799F366 Mg-sulphate.95	0.039407	0.049956	0.045877	0.049218
799F366 Mg-sulphate.96	0.047049	0.046195	0.050887	0.058012
799F366 Mg-sulphate.97	0.043125	0.050171	0.049397	0.054158
799F366 Mg-sulphate.98	0.045607	0.051494	0.043519	0.052846
799F366 Mg-sulphate.99	0.054018	0.048240	0.045879	0.056317

3070.29515395 3080.04791767 3089.80344009 \

Data ID			
CAGR01 Calcite.	0.696149	0.585477	0.687150
CAGR01 Calcite.1	0.673004	0.651316	0.561063
CAGR01 Calcite.10	0.551777	0.632937	0.574604
CAGR01 Calcite.100	0.566445	0.540179	0.606870

CAGR01 Calcite.101	0.595441	0.714247	0.619907
...	...	...	...
799F366 Mg-sulphate.95	0.052789	0.050889	0.054148
799F366 Mg-sulphate.96	0.048558	0.058945	0.045808
799F366 Mg-sulphate.97	0.051973	0.051844	0.057178
799F366 Mg-sulphate.98	0.057114	0.054627	0.055296
799F366 Mg-sulphate.99	0.060730	0.062376	0.047758

3099.5617061

Data ID

CAGR01 Calcite.	0.659407
CAGR01 Calcite.1	0.643262
CAGR01 Calcite.10	0.666354
CAGR01 Calcite.100	0.596159
CAGR01 Calcite.101	0.594143

...	...
799F366 Mg-sulphate.95	0.063786
799F366 Mg-sulphate.96	0.057465
799F366 Mg-sulphate.97	0.052170
799F366 Mg-sulphate.98	0.058471
799F366 Mg-sulphate.99	0.060878

[3712 rows x 316 columns]

```
[52]: axes = obs.plot_profiles(scope='all', groupby='Category', stacked=False,
    ↪waterfall=True, ci=False)
```

Enfys_Inas_Mwir Sampled All Entries by Category with Noise




Figure 15 Waterfall diagram of the entries of the MICA Files, as sampled by *Enfys*, including synthetic noise. Each entry bar is composed of n repeat noisy samples, such that each bar is composed of n rows, illustrating the noise.

We need to recompute the continuum-removed data on this noisy set.

```
[53]: cr_tool = sla(obs)
      cr_obs = cr_tool.remove_continuum()

[54]: axes = cr_obs.plot_profiles(scope='all', groupby='Category', stacked=False,
      ↪waterfall=True, ci=False)
```

Enfys_Inas_Mwir Sampled All Entries by Category
Continuum Removed with Noise






Figure 16 Continuum-Removed waterfall diagram of the entries of the MICA Files, as sampled by *Enfys*, including synthetic noise. Each entry bar is composed of n repeat noisy samples, such that each bar is composed of n rows, illustrating the noise.

```
[55]: axes = cr_obs.plot_profiles(scope='samples', groupby='Species', stacked=False,
    ↪waterfall=True, ci=False)
```


Enfys_Inas_Mwir Sampled Samples grouped by Species
Continuum Removed with Noise (1/5)




Enfys_Inas_Mwir Sampled Samples grouped by Species
Continuum Removed with Noise (2/5)



Enfys_Inas_Mwir Sampled Samples grouped by Species
Continuum Removed with Noise (3/5)



Enfys_Inas_Mwir Sampled Samples grouped by Species
Continuum Removed with Noise (4/5)



Enfys_Inas_Mwir Sampled Samples grouped by Species
Continuum Removed with Noise (5/5)




Figure 17 Detailed waterfall view of each entry in the spectral library, showing the scale of noise

introduced, relative to the features of each spectrum, and the reflectance range. Each entry bar is composed of n repeat noisy samples, such that each bar is composed of n rows, illustrating the noise.

5.5.1 Validating the Noise

We can check that our noise is correct by aggregating the repeat samples to give the mean and standard deviation for each entry and each spectral channel (wavelength).

We can then compare this to the spectral SNR.

```
[56]: signal = obs.main_df.groupby('Root Data ID').mean()
      noise = obs.main_df.groupby('Root Data ID').std()
      snr = signal/noise
```

```
/var/folders/w4/44k32f413dd82lxmtw9456w80000gp/T/ipykernel_76523/2651352844.py:1
: FutureWarning: The default value of numeric_only in DataFrameGroupBy.mean is
deprecated. In a future version, numeric_only will default to False. Either
specify numeric_only or select only columns which should be valid for the
function.
```

```
      signal = obs.main_df.groupby('Root Data ID').mean()
/var/folders/w4/44k32f413dd82lxmtw9456w80000gp/T/ipykernel_76523/2651352844.py:2
: FutureWarning: The default value of numeric_only in DataFrameGroupBy.std is
deprecated. In a future version, numeric_only will default to False. Either
specify numeric_only or select only columns which should be valid for the
function.
      noise = obs.main_df.groupby('Root Data ID').std()
```

```
[57]: snr.mean()
```


```
[57]: 900.000000      18995.653023
      903.285743      19300.937699
      906.618278      19041.784707
      909.999291      19439.110070
      913.430570      19257.049290
      ...
      3060.545164      12.542792
      3070.295154      12.580475
      3080.047918      12.428545
      3089.803440      12.405591
      3099.561706      12.136643
      Length: 299, dtype: float64
```

```
[58]: # plot the SNR
      import seaborn as sns
      # plot with log scale y axis
      g = sns.lineplot(data=snr.mean(), label='Mean/Standard Deviation of Sampled_
      ↪Noisy Data', markers=True)
      g.set(yscale="log")
```

```
# get instrument snr
sns.lineplot(data=obs.instrument.main_df, x='cwl', y='snr', label='InGaAs + InAs SNR', linestyle='--', markers=True)

g.set_xlabel('Wavelength (nm)')
g.set_ylabel('SNR')
```

[58]: Text(0, 0.5, 'SNR')



5.6 Spectral Angle Mapper

TODO: perform spectral angle mapper evaluation of each spectrum against the other spectra of the set.

6 Discussion

Notes on what we found above, perhaps with some more analysis.

7 Conclusions

Concluding remarks and summary of findings.

8 References